

# Problem Statement

To estimate the delivery time of Porter using Neural Networks

## Importing libraries

```
In [1]: # Importing necessary libraries for data manipulation and analysis
import pandas as pd
import numpy as np

# Importing libraries for data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Importing libraries for handling datetime operations
from datetime import datetime

# Importing libraries for preprocessing and encoding
from sklearn.preprocessing import StandardScaler, OneHotEncoder

# Importing libraries for splitting data
from sklearn.model_selection import train_test_split

# Importing libraries for neural network
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping

# Importing libraries for evaluation metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error

# Setting up visualization styles
sns.set(style='whitegrid')
plt.rcParams['figure.figsize'] = (10, 6)

# Ignoring warnings
import warnings
warnings.filterwarnings('ignore')
```

## Importing Porter data

```
In [2]: df=pd.read_csv('porter_data.csv')
df.head()
```

Out[2]:

	market_id	created_at	actual_delivery_time	store_id	store_primary_c
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	

## Printing basic info of the features

In [3]:

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            196441 non-null float64
1   created_at                           197428 non-null object
2   actual_delivery_time                  197421 non-null object
3   store_id                             197428 non-null object
4   store_primary_category                192668 non-null object
5   order_protocol                       196433 non-null float64
6   total_items                          197428 non-null int64
7   subtotal                             197428 non-null int64
8   num_distinct_items                   197428 non-null int64
9   min_item_price                       197428 non-null int64
10  max_item_price                       197428 non-null int64
11  total_onshift_partners                181166 non-null float64
12  total_busy_partners                  181166 non-null float64
13  total_outstanding_orders              181166 non-null float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB
```

## Converting relevant columns to pandas Datetime

In [4]:

```
# Converting 'created_at' and 'actual_delivery_time' to datetime format
df['created_at'] = pd.to_datetime(df['created_at'])
df['actual_delivery_time'] = pd.to_datetime(df['actual_delivery_time'])
```

In [5]:

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            196441 non-null  float64
1   created_at                           197428 non-null  datetime64[ns]
2   actual_delivery_time                 197421 non-null  datetime64[ns]
3   store_id                             197428 non-null  object
4   store_primary_category               192668 non-null  object
5   order_protocol                       196433 non-null  float64
6   total_items                          197428 non-null  int64
7   subtotal                             197428 non-null  int64
8   num_distinct_items                  197428 non-null  int64
9   min_item_price                       197428 non-null  int64
10  max_item_price                       197428 non-null  int64
11  total_onshift_partners                181166 non-null  float64
12  total_busy_partners                   181166 non-null  float64
13  total_outstanding_orders              181166 non-null  float64
dtypes: datetime64[ns](2), float64(5), int64(5), object(2)
memory usage: 21.1+ MB

```

## Creating target column(Time taken)

```

In [13]: # Create a new column named 'time_taken' to store the difference in minutes
df['time_taken'] = (df['actual_delivery_time'] - df['created_at'])

```

```

In [14]: df.head()

```

```

Out[14]:
   market_id  created_at  actual_delivery_time  store_id  store_primary_c
0          1.0  2015-02-06 2015-02-06 23:27:16  df263d996281d984952c07998dc54358
1          2.0  2015-02-10 2015-02-10 22:56:29  f0ade77b43923b38237db569b016ba25
2          3.0  2015-01-22 2015-01-22 21:09:09  f0ade77b43923b38237db569b016ba25
3          3.0  2015-02-03 2015-02-03 22:13:00  f0ade77b43923b38237db569b016ba25
4          3.0  2015-02-15 2015-02-15 03:20:26  f0ade77b43923b38237db569b016ba25

```

```

In [15]: # Extracting the total minutes from the 'time_taken' column
df['time_taken_minutes'] = df['time_taken'].dt.total_seconds() // 60

```

```

In [16]: df.head()

```

Out[16]:

	market_id	created_at	actual_delivery_time	store_id	store_primary_c
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	ε
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	

## Feature Engineering and Data Preprocessing

### Creating hour and day of week column

In [17]:

```
# Extracting hour and day of the week from 'created_at'
df['order_hour'] = df['created_at'].dt.hour
df['order_day_of_week'] = df['created_at'].dt.dayofweek # Monday=0, Sunday=6
```

In [18]:

```
df.head()
```

Out[18]:

	market_id	created_at	actual_delivery_time	store_id	store_primary_c
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	ε
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	

### Dropping columns that arent useful anymore

In [19]:

```
df.drop(['time_taken', 'created_at', 'actual_delivery_time'], axis=1, inplace=True)
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   market_id                            196441 non-null  float64
1   store_id                             197428 non-null  object
2   store_primary_category               192668 non-null  object
3   order_protocol                       196433 non-null  float64
4   total_items                          197428 non-null  int64
5   subtotal                             197428 non-null  int64
6   num_distinct_items                   197428 non-null  int64
7   min_item_price                       197428 non-null  int64
8   max_item_price                       197428 non-null  int64
9   total_onshift_partners                181166 non-null  float64
10  total_busy_partners                  181166 non-null  float64
11  total_outstanding_orders              181166 non-null  float64
12  time_taken_minutes                   197421 non-null  float64
13  order_hour                           197428 non-null  int32
14  order_day_of_week                    197428 non-null  int32
dtypes: float64(6), int32(2), int64(5), object(2)
memory usage: 21.1+ MB

```

In [ ]:

## Handling Null values

In [20]: `df.isna().sum()`

```

Out[20]: market_id                987
store_id                0
store_primary_category  4760
order_protocol          995
total_items              0
subtotal                0
num_distinct_items      0
min_item_price          0
max_item_price          0
total_onshift_partners  16262
total_busy_partners     16262
total_outstanding_orders 16262
time_taken_minutes      7
order_hour              0
order_day_of_week       0
dtype: int64

```

```

In [25]: # Finding the number of unique values in each column
unique_values = {column: df[column].nunique() for column in df.columns}

# Displaying the unique values count for each column
for column, unique_count in unique_values.items():
    print(f"{column}: {unique_count}")

```

```

market_id: 6
store_id: 6743
store_primary_category: 74
order_protocol: 7
total_items: 57
subtotal: 8368
num_distinct_items: 20
min_item_price: 2312
max_item_price: 2652
total_onshift_partners: 172
total_busy_partners: 159
total_outstanding_orders: 281
time_taken_minutes: 274
order_hour: 19
order_day_of_week: 7

```

```
In [71]: df1=df.dropna()
```

```
In [93]: df[df["store_id"]=="252a3dbaeb32e7690242ad3b556e626b"]
```

```
Out[93]:
```

	market_id	store_id	store_primary_category	order_protocol	tot
<b>52018</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
<b>52019</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
<b>52020</b>	2.0	252a3dbaeb32e7690242ad3b556e626b	burger	3.0	
<b>52021</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
<b>52022</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
...	...	...	...	...	...
<b>63432</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
<b>63433</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
<b>63434</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
<b>63435</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	
<b>63436</b>	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	

350 rows × 15 columns

## Checking whether mean or median is the right choice for Null imputation

```
In [100]: df.groupby("market_id")["total_onshift_partners"].mean()
```

```
Out[100]:
market_id
1.0      24.208854
2.0      62.590695
3.0      18.847580
4.0      60.464482
5.0      23.911045
6.0      44.929771
Name: total_onshift_partners, dtype: float64
```

```
In [101]: df.groupby("market_id")["total_onshift_partners"].median()
```

```
Out[101]: market_id
1.0      19.0
2.0      55.0
3.0      15.0
4.0      60.0
5.0      20.0
6.0      36.0
Name: total_onshift_partners, dtype: float64
```

```
In [103]: df.groupby("order_hour")["total_onshift_partners"].mean()
```

```
Out[103]: order_hour
0      27.933751
1      54.325601
2      67.995169
3      64.205588
4      44.996112
5      23.589613
6      13.421094
7      10.777778
8       0.000000
14     0.550000
15     2.141473
16     4.965949
17     7.757729
18    15.092275
19    32.199487
20    37.353387
21    30.325540
22    22.749043
23    20.274580
Name: total_onshift_partners, dtype: float64
```

```
In [105]: df.groupby("order_day_of_week")["total_onshift_partners"].mean()
```

```
Out[105]: order_day_of_week
0      42.084044
1      37.333062
2      40.067352
3      43.746503
4      48.602855
5      52.111917
6      45.943654
Name: total_onshift_partners, dtype: float64
```

```
In [112]: df.groupby(["market_id", "order_hour"])["total_onshift_partners"].mean()
```

```
Out[112]: market_id order_hour
1.0      0      14.437811
         1      26.014145
         2      36.809734
         3      37.072227
         4      27.385254
         ...
6.0     19      30.744186
         20      40.627907
         21      31.200000
         22      23.806452
         23      18.000000
Name: total_onshift_partners, Length: 106, dtype: float64
```

## Mean Imputation

```
In [113... # List of columns to impute
columns_to_impute = ['total_outstanding_orders', 'total_busy_partners', 'total_onsh

# Group by 'market_id' and 'order_hour'
grouped = df.groupby(['market_id', 'order_hour'])

# Impute missing values
for column in columns_to_impute:
    # Calculate the mean for each group and transform to align with the original Da
    df[column] = grouped[column].transform(lambda x: x.fillna(x.mean()))
```

In [114... df

```
Out[114]:
```

	market_id	store_id	store_primary_category	order_protocol	tc
0	1.0	df263d996281d984952c07998dc54358	american	1.0	
1	2.0	f0ade77b43923b38237db569b016ba25	mexican	2.0	
2	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
3	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
4	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
...	...	...	...	...	...
197423	1.0	a914ecef9c12ffdb9bede64bb703d877	fast	4.0	
197424	1.0	a914ecef9c12ffdb9bede64bb703d877	fast	4.0	
197425	1.0	a914ecef9c12ffdb9bede64bb703d877	fast	4.0	
197426	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	
197427	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	

197428 rows × 15 columns

```
In [115... df.isna().sum()
```

```
Out[115]:
```

market_id	987
store_id	0
store_primary_category	4760
order_protocol	995
total_items	0
subtotal	0
num_distinct_items	0
min_item_price	0
max_item_price	0
total_onshift_partners	989
total_busy_partners	989
total_outstanding_orders	989
time_taken_minutes	7
order_hour	0
order_day_of_week	0

dtype: int64

## Dropping null rows

```
In [117... df[df["total_onshift_partners"].isnull()].dropna(inplace=True)
```



```
In [118... df.isna().sum()
```

```
Out[118]: market_id          987
store_id            0
store_primary_category  4760
order_protocol      995
total_items         0
subtotal            0
num_distinct_items   0
min_item_price       0
max_item_price       0
total_onshift_partners  989
total_busy_partners   989
total_outstanding_orders  989
time_taken_minutes    7
order_hour           0
order_day_of_week     0
dtype: int64
```

```
In [119... df= df[~df['total_onshift_partners'].isnull()]
```

```
In [121... df.isna().sum()
```

```
Out[121]: market_id          0
store_id            0
store_primary_category  4268
order_protocol      508
total_items         0
subtotal            0
num_distinct_items   0
min_item_price       0
max_item_price       0
total_onshift_partners  0
total_busy_partners   0
total_outstanding_orders  0
time_taken_minutes    7
order_hour           0
order_day_of_week     0
dtype: int64
```

```
In [122... df= df[~df['order_protocol'].isnull()]
```

```
In [123... df.isna().sum()
```

```
Out[123]: market_id          0
store_id            0
store_primary_category  4005
order_protocol        0
total_items          0
subtotal             0
num_distinct_items   0
min_item_price       0
max_item_price       0
total_onshift_partners  0
total_busy_partners   0
total_outstanding_orders  0
time_taken_minutes    7
order_hour           0
order_day_of_week     0
dtype: int64
```

```
In [124... df= df[~df['time_taken_minutes'].isnull()]
```

In [125... `df.isna().sum()`

```
Out[125]: market_id          0
store_id          0
store_primary_category  4005
order_protocol     0
total_items        0
subtotal           0
num_distinct_items  0
min_item_price     0
max_item_price     0
total_onshift_partners  0
total_busy_partners  0
total_outstanding_orders  0
time_taken_minutes  0
order_hour         0
order_day_of_week   0
dtype: int64
```

In [126... `df[df["store_primary_category"].isna()]`

```
Out[126]:
```

	market_id	store_id	store_primary_category	order_protocol	tc
2	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
3	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
4	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
5	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
6	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
...	...	...	...	...	...
197208	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	
197209	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	
197210	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	
197211	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	
197212	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	

4005 rows × 15 columns

In [128... `df[df["store_id"]=="f0ade77b43923b38237db569b016ba25"]`

Out[128]:	market_id	store_id	store_primary_category	order_protocol	total_it
1	2.0	f0ade77b43923b38237db569b016ba25	mexican	2.0	
2	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
3	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
4	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
5	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
6	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
7	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
8	2.0	f0ade77b43923b38237db569b016ba25	indian	3.0	
9	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
10	3.0	f0ade77b43923b38237db569b016ba25	NaN	4.0	
11	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
12	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	
13	3.0	f0ade77b43923b38237db569b016ba25	NaN	4.0	

```
In [129... df[df["store_primary_category"].isna()]["store_id"].nunique()
```

Out[129]: 632

```
In [131... df2=df[df["store_primary_category"].isna()]["store_id"].unique()
```

## Imputing store\_primary\_category by mode

```
In [138... # Function to impute missing values by mode, handling ties randomly
def impute_by_mode(df, column):
    # Get the mode(s)
    modes = df[column].mode()

    if len(modes) > 1:
        # If there are ties, choose one randomly with equal probability
        chosen_mode = np.random.choice(modes)
    else:
        # If no tie, use the single mode
        chosen_mode = modes[0]

    # Impute missing values with the chosen mode
    df[column].fillna(chosen_mode, inplace=True)

# List of columns to impute
columns_to_impute = ['store_primary_category']

# Apply the function to each column
for column in columns_to_impute:
    impute_by_mode(df, column)
```

```
In [139... df.isna().sum()
```

```
Out[139]: market_id      0
store_id      0
store_primary_category  0
order_protocol  0
total_items    0
subtotal       0
num_distinct_items  0
min_item_price  0
max_item_price  0
total_onshift_partners  0
total_busy_partners  0
total_outstanding_orders  0
time_taken_minutes  0
order_hour      0
order_day_of_week  0
dtype: int64
```

```
In [140... df.shape
```

```
Out[140]: (195924, 15)
```

```
In [142... df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 195924 entries, 0 to 197427
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            195924 non-null float64
1   store_id                             195924 non-null object
2   store_primary_category               195924 non-null object
3   order_protocol                       195924 non-null float64
4   total_items                          195924 non-null int64
5   subtotal                            195924 non-null int64
6   num_distinct_items                  195924 non-null int64
7   min_item_price                      195924 non-null int64
8   max_item_price                      195924 non-null int64
9   total_onshift_partners              195924 non-null float64
10  total_busy_partners                 195924 non-null float64
11  total_outstanding_orders            195924 non-null float64
12  time_taken_minutes                  195924 non-null float64
13  order_hour                          195924 non-null int32
14  order_day_of_week                   195924 non-null int32
dtypes: float64(6), int32(2), int64(5), object(2)
memory usage: 22.4+ MB
```

```
In [143... store_name_counts = df['store_id'].value_counts()
df['store_name_enc'] = df['store_id'].map(store_name_counts)
```

```
In [148... df = df.drop('store_name_enc', axis=1)
```

```
In [149... df
```

Out[149]:

		market_id	store_id	store_primary_category	order_protocol	tc
0	1.0	df263d996281d984952c07998dc54358		american	1.0	
1	2.0	f0ade77b43923b38237db569b016ba25		mexican	2.0	
2	3.0	f0ade77b43923b38237db569b016ba25		american	1.0	
3	3.0	f0ade77b43923b38237db569b016ba25		american	1.0	
4	3.0	f0ade77b43923b38237db569b016ba25		american	1.0	
...	...		...	...	...	...
197423	1.0	a914ecef9c12ffdb9bede64bb703d877		fast	4.0	
197424	1.0	a914ecef9c12ffdb9bede64bb703d877		fast	4.0	
197425	1.0	a914ecef9c12ffdb9bede64bb703d877		fast	4.0	
197426	1.0	c81e155d85dae5430a8cee6f2242e82c		sandwich	1.0	
197427	1.0	c81e155d85dae5430a8cee6f2242e82c		sandwich	1.0	

195924 rows × 15 columns

## Using Label Encoding for store name

In [150... `from sklearn.preprocessing import LabelEncoder`

In [152... `label_encoder = LabelEncoder()`  
`df['store_name_encoded'] = label_encoder.fit_transform(df['store_id'])`

In [153... `df`

Out[153]:

		market_id	store_id	store_primary_category	order_protocol	tc
0	1.0	df263d996281d984952c07998dc54358		american	1.0	
1	2.0	f0ade77b43923b38237db569b016ba25		mexican	2.0	
2	3.0	f0ade77b43923b38237db569b016ba25		american	1.0	
3	3.0	f0ade77b43923b38237db569b016ba25		american	1.0	
4	3.0	f0ade77b43923b38237db569b016ba25		american	1.0	
...	...		...	...	...	...
197423	1.0	a914ecef9c12ffdb9bede64bb703d877		fast	4.0	
197424	1.0	a914ecef9c12ffdb9bede64bb703d877		fast	4.0	
197425	1.0	a914ecef9c12ffdb9bede64bb703d877		fast	4.0	
197426	1.0	c81e155d85dae5430a8cee6f2242e82c		sandwich	1.0	
197427	1.0	c81e155d85dae5430a8cee6f2242e82c		sandwich	1.0	

195924 rows × 16 columns

In [154... `df=df.drop("store_id",axis=1)`

In [155...

df

Out[155]:

	market_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_iter
0	1.0	american	1.0	4	3441	
1	2.0	mexican	2.0	1	1900	
2	3.0	american	1.0	1	1900	
3	3.0	american	1.0	6	6900	
4	3.0	american	1.0	3	3900	
...	...	...	...	...	...	...
197423	1.0	fast	4.0	3	1389	
197424	1.0	fast	4.0	6	3010	
197425	1.0	fast	4.0	5	1836	
197426	1.0	sandwich	1.0	1	1175	
197427	1.0	sandwich	1.0	4	2605	

195924 rows × 15 columns

In [156...

```
duplicates = df.duplicated()
```

```
# Print the original DataFrame with a marker for duplicates
print(df.loc[duplicates])
```

```

      market_id store_primary_category order_protocol total_items \
139263        6.0                indian            3.0          2
166281        6.0                cafe            4.0          1

      subtotal num_distinct_items min_item_price max_item_price \
139263      1650                1           825           825
166281       350                1           350           350

      total_onshift_partners total_busy_partners total_outstanding_orders \
139263          39.813559          40.40678          51.135593
166281          39.813559          40.40678          51.135593

      time_taken_minutes order_hour order_day_of_week store_name_encoded
139263             24.0          4             1          2637
166281             39.0          4             4          1501

```

In [159...


```
df=df.drop_duplicates()
```

In [160...

df

Out[160]:	market_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_iter
	0	1.0	american	1.0	4	3441
	1	2.0	mexican	2.0	1	1900
	2	3.0	american	1.0	1	1900
	3	3.0	american	1.0	6	6900
	4	3.0	american	1.0	3	3900
	...	...	...	...	...	...
	197423	1.0	fast	4.0	3	1389
	197424	1.0	fast	4.0	6	3010
	197425	1.0	fast	4.0	5	1836
	197426	1.0	sandwich	1.0	1	1175
	197427	1.0	sandwich	1.0	4	2605

195922 rows × 15 columns



In [161... `df.isna().sum()`

Out[161]:

market_id	0
store_primary_category	0
order_protocol	0
total_items	0
subtotal	0
num_distinct_items	0
min_item_price	0
max_item_price	0
total_onshift_partners	0
total_busy_partners	0
total_outstanding_orders	0
time_taken_minutes	0
order_hour	0
order_day_of_week	0
store_name_encoded	0
dtype:	int64

In [163... `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 195922 entries, 0 to 197427
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   market_id                            195922 non-null float64
1   store_primary_category              195922 non-null object
2   order_protocol                      195922 non-null float64
3   total_items                         195922 non-null int64
4   subtotal                           195922 non-null int64
5   num_distinct_items                 195922 non-null int64
6   min_item_price                     195922 non-null int64
7   max_item_price                     195922 non-null int64
8   total_onshift_partners              195922 non-null float64
9   total_busy_partners                 195922 non-null float64
10  total_outstanding_orders            195922 non-null float64
11  time_taken_minutes                  195922 non-null float64
12  order_hour                          195922 non-null int32
13  order_day_of_week                   195922 non-null int32
14  store_name_encoded                  195922 non-null int32
dtypes: float64(6), int32(3), int64(5), object(1)
memory usage: 21.7+ MB
```

## label Encoding store\_primary\_category

```
In [164... label_encoder = LabelEncoder()
df['store_primary_category_enc'] = label_encoder.fit_transform(df['store_primary_ca
```

```
In [166... df=df.drop("store_primary_category",axis=1)
```

```
In [ ]:
```

```
In [167... df
```

```
Out[167]:
```

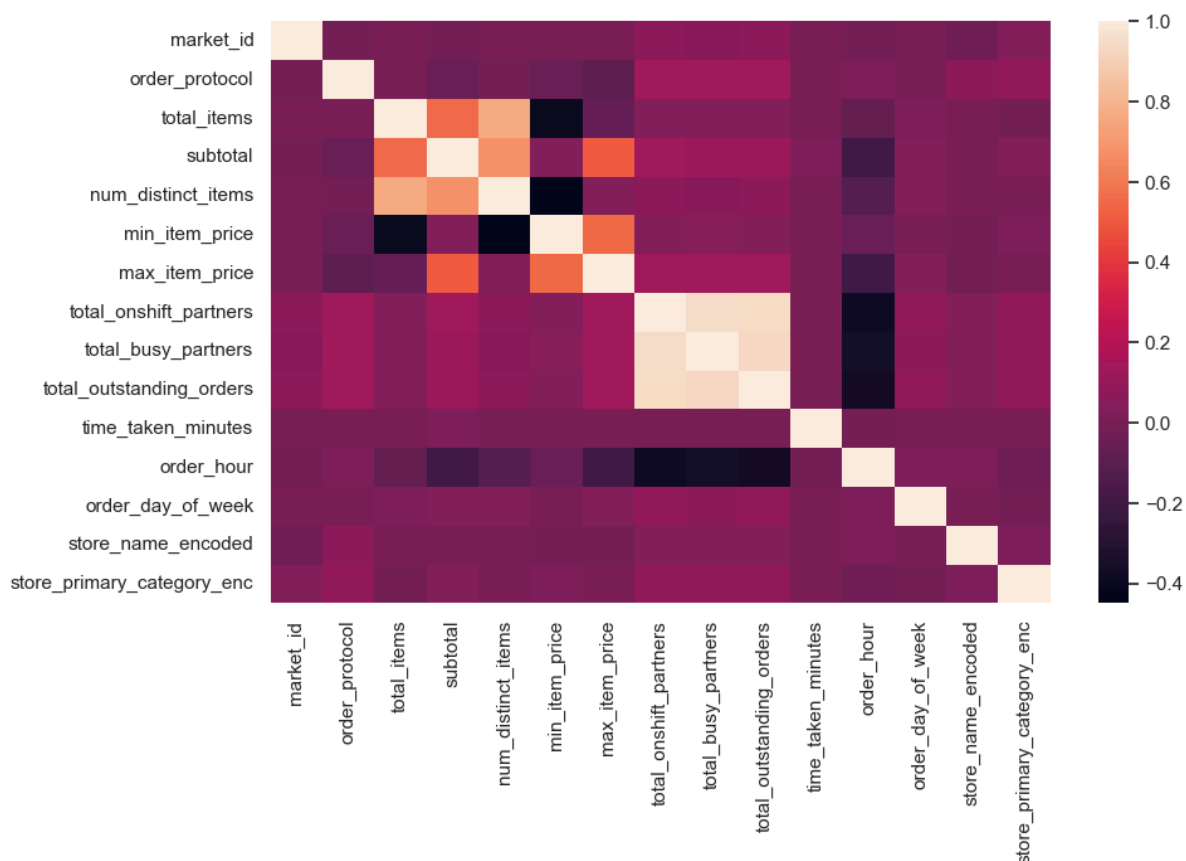
	market_id	order_protocol	total_items	subtotal	num_distinct_items	min_item_price	max
0	1.0	1.0	4	3441	4	557	
1	2.0	2.0	1	1900	1	1400	
2	3.0	1.0	1	1900	1	1900	
3	3.0	1.0	6	6900	5	600	
4	3.0	1.0	3	3900	3	1100	
...	...	...	...	...	...	...	...
197423	1.0	4.0	3	1389	3	345	
197424	1.0	4.0	6	3010	4	405	
197425	1.0	4.0	5	1836	3	300	
197426	1.0	1.0	1	1175	1	535	
197427	1.0	1.0	4	2605	4	425	

195922 rows × 15 columns



```
In [168...] sns.heatmap(df.corr())
```

```
Out[168]: <Axes: >
```




```
In [169...] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 195922 entries, 0 to 197427
Data columns (total 15 columns):
 #   Column                                  Non-Null Count  Dtype  
---  -
 0   market_id                             195922 non-null float64
 1   order_protocol                         195922 non-null float64
 2   total_items                           195922 non-null int64
 3   subtotal                              195922 non-null int64
 4   num_distinct_items                    195922 non-null int64
 5   min_item_price                        195922 non-null int64
 6   max_item_price                        195922 non-null int64
 7   total_onshift_partners                 195922 non-null float64
 8   total_busy_partners                   195922 non-null float64
 9   total_outstanding_orders               195922 non-null float64
10   time_taken_minutes                     195922 non-null float64
11   order_hour                            195922 non-null int32
12   order_day_of_week                     195922 non-null int32
13   store_name_encoded                    195922 non-null int32
14   store_primary_category_enc             195922 non-null int32
dtypes: float64(6), int32(4), int64(5)
memory usage: 20.9 MB
```


```
In [178...] # Create the scatter plot
sns.scatterplot(x='time_taken_minutes', y='subtotal', data=df)

# Set the x-axis limit
plt.xlim(0, 800)
```

```
# Show the plot
plt.show()
```



```
In [173]: sns.scatterplot(x='total_items', y='subtotal', hue='num_distinct_items', palette='cividis')
Out[173]: <Axes: xlabel='total_items', ylabel='subtotal'>
```



```
In [222]: df3=df.copy()
```

```
In [223]: df3.shape
```

```
Out[223]: (195922, 15)
```

```
In [225]: df3=df3.drop("store_name_encoded",axis=1)
```

## Removing outliers using LOF

```
In [228... from sklearn.neighbors import LocalOutlierFactor
import matplotlib.pyplot as plt
model1 = LocalOutlierFactor(contamination=0.05)
df3['lof_anomaly_score'] = model1.fit_predict(df3)
```


```
In [229... print("number of outliers : ",(len(df3.loc[(df3['lof_anomaly_score'] == -1)])))
df3=df3.loc[(df3['lof_anomaly_score'] == 1)]

number of outliers : 8817
```

```
In [230... df3.drop(['lof_anomaly_score'],axis=1,inplace=True)
```

```
In [231... # Create the scatter plot
sns.scatterplot(x='time_taken_minutes', y='subtotal', data=df3)
```

```
Out[231]: <Axes: xlabel='time_taken_minutes', ylabel='subtotal'>
```

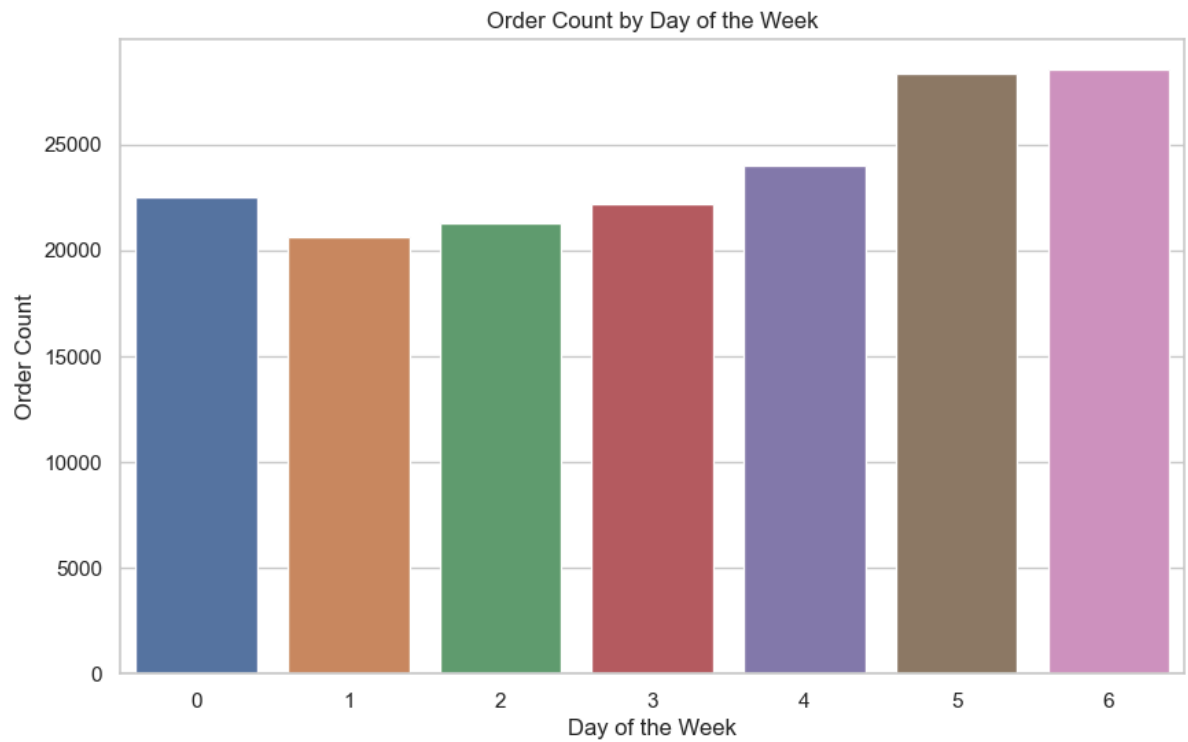


## Making various plots from features

```
In [232... # Create a countplot for the 'order_day_of_week' column
sns.countplot(x='order_day_of_week', data=df3)

# Set the title and labels
plt.title('Order Count by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Order Count')


# Show the plot
plt.show()
```




```
In [233...] # Create a countplot for the 'order_day_of_week' column
sns.countplot(x='order_hour', data=df3)

# Set the title and labels
plt.title('Order Count by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Order Count')

# Show the plot
plt.show()
```




```
In [276...] sns.countplot(x=df.market_id)
plt.show()
```



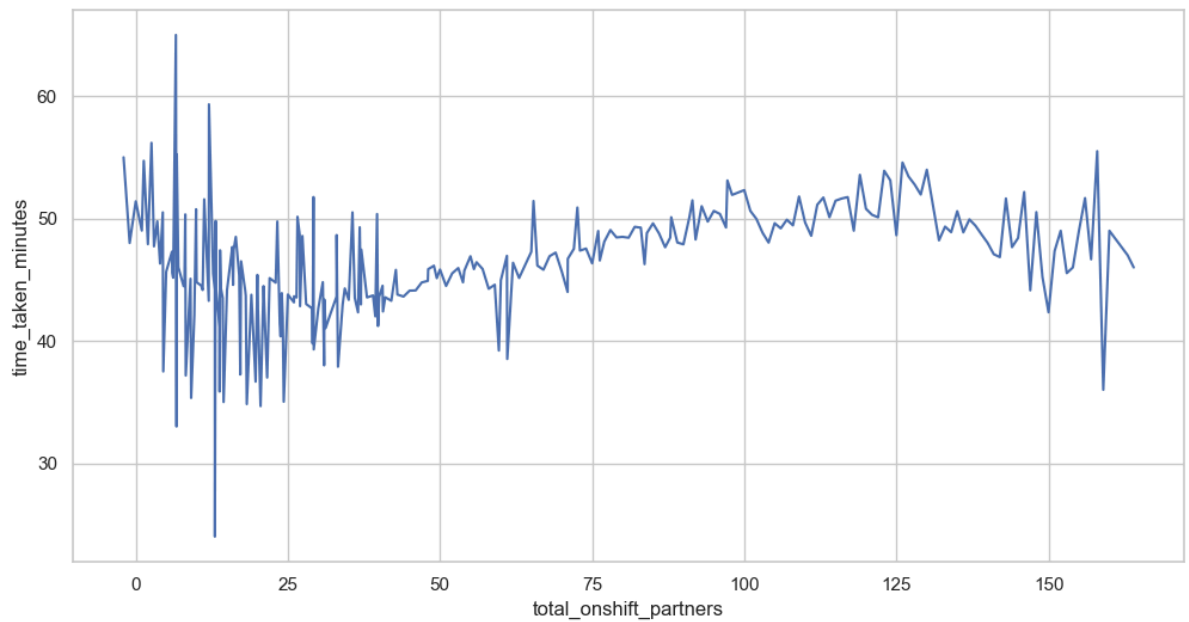
```
In [234... # Create a scatter plot for 'order_hour' vs 'time_taken_minutes'
plt.figure(figsize=(12, 6))
sns.scatterplot(x='order_hour', y='time_taken_minutes', data=df3)

# Set the title and labels
plt.title('Time Taken Minutes vs Order Hour')
plt.xlabel('Order Hour')
plt.ylabel('Time Taken Minutes')

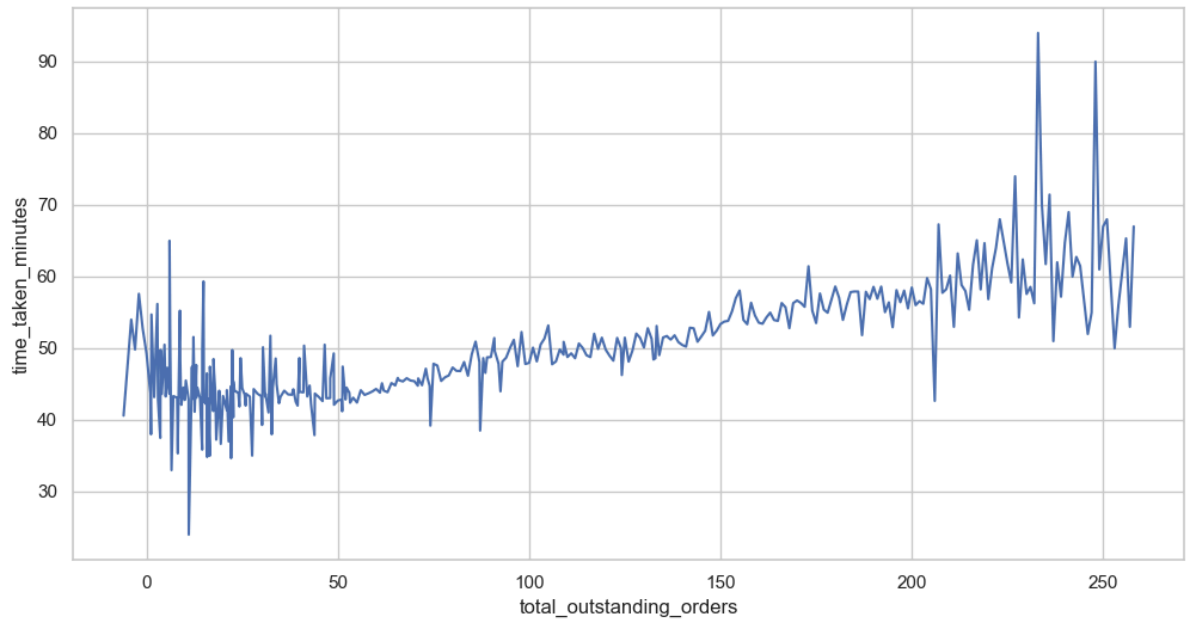
# Show the plot
plt.show()
```




```
In [277... plt.figure(figsize=(12, 6))
sns.lineplot(x='total_onshift_partners', y='time_taken_minutes', data=df3, ci=None)
plt.show()
```



```
In [278... plt.figure(figsize=(12, 6))
sns.lineplot(x='total_outstanding_orders', y='time_taken_minutes', data=df3, ci=None)
plt.show()
```




```
In [279... plt.figure(figsize=(12, 6))
sns.lineplot(x='total_busy_partners', y='time_taken_minutes', data=df3, ci=None)
plt.show()
```



```
In [235... plt.figure(figsize=(12, 6))
sns.lineplot(x='order_hour', y='time_taken_minutes', data=df3, ci=None)

# Set the title and labels
plt.title('Average Time Taken Minutes vs Order Hour')
plt.xlabel('Order Hour')
plt.ylabel('Time Taken Minutes')


# Show the plot
plt.show()
```



```
In [236... # Create a box plot for 'order_hour' vs 'time_taken_minutes'
plt.figure(figsize=(12, 6))
sns.boxplot(x='order_hour', y='time_taken_minutes', data=df3)

# Set the title and labels
plt.title('Distribution of Time Taken Minutes vs Order Hour')
plt.xlabel('Order Hour')
plt.ylabel('Time Taken Minutes')

# Show the plot
plt.show()
```



```
In [237... y=df3['time_taken_minutes']
x = df3.drop(['time_taken_minutes'], axis=1)

X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_st
```

```
In [267... x
```

Out[267]:

	market_id	order_protocol	total_items	subtotal	num_distinct_items	min_item_price	max
<b>0</b>	1.0	1.0	4	3441	4	557	
<b>1</b>	2.0	2.0	1	1900	1	1400	
<b>2</b>	3.0	1.0	1	1900	1	1900	
<b>3</b>	3.0	1.0	6	6900	5	600	
<b>4</b>	3.0	1.0	3	3900	3	1100	
...	...	...	...	...	...	...	...
<b>197422</b>	1.0	4.0	7	2445	3	145	
<b>197423</b>	1.0	4.0	3	1389	3	345	
<b>197424</b>	1.0	4.0	6	3010	4	405	
<b>197425</b>	1.0	4.0	5	1836	3	300	
<b>197427</b>	1.0	1.0	4	2605	4	425	

167512 rows × 13 columns

```
In [239... y
```



```
Out[239]: 0          62.0
          1          67.0
          2          29.0
          3          51.0
          4          39.0
          ...
          197422      39.0
          197423      65.0
          197424      56.0
          197425      50.0
          197427      37.0
          Name: time_taken_minutes, Length: 167512, dtype: float64
```

```
In [240... #random forest model training
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.metrics import mean_absolute_error
from sklearn.ensemble import RandomForestRegressor
```

## Creating baseline model RF to compare with Neural Networks

```
In [241... regressor = RandomForestRegressor()

regressor.fit(X_train, y_train)
```

```
Out[241]: ▼ RandomForestRegressor
          RandomForestRegressor()
```

```
In [242... prediction = regressor.predict(X_test)
mse = mean_squared_error(y_test, prediction)
rmse = mse**.5
print("mse : ", mse)
print("rmse : ", rmse)
mae = mean_absolute_error(y_test, prediction)
print('mae:' ,mae)
```

```
mse : 189.7520443763879
rmse : 13.775051519917735
mae: 10.596033335678541
```

```
In [243... r2_score(y_test, prediction)
```

```
Out[243]: 0.2674632758036619
```


```
In [210... def MAPE(Y_actual,Y_Predicted):
    mape = np.mean(np.abs((Y_actual - Y_Predicted)/Y_actual))*100
    return mape
```

```
In [211... print("mape : ",MAPE(y_test, prediction))

mape : 26.785473813844728
```

```
In [212... sorted_idx = regressor.feature_importances_.argsort()
plt.barh(df3.columns[sorted_idx], regressor.feature_importances_[sorted_idx])
plt.xlabel("Random Forest Feature Importance")
```

```
Out[212]: Text(0.5, 0, 'Random Forest Feature Importance')
```



## Train-Test Splitting Standard Scaling

```
In [268... from sklearn import preprocessing
scaler = preprocessing.MinMaxScaler()
x_scaled = scaler.fit_transform(x)
X_train, X_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2, ra
```


## Creating Neural Network Architecture

```
In [269... model = Sequential()
model.add(Dense(11, kernel_initializer='normal'))
model.add(Dense(256, activation='relu'))
model.add(Dense(512, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(1, activation='linear'))
```


## Model Training

```
In [270... from tensorflow.keras.optimizers import Adam
adam=Adam(learning_rate=0.01)
model.compile(loss='mse', optimizer=adam, metrics=['mse', 'mae'])
history=model.fit(X_train, y_train, epochs=30, batch_size=512, verbose=1, validatio
```


Epoch 1/30

**210/210**  5s 16ms/step - loss: 494.4725 - mae: 16.5599 - mse: 494.4740 - val\_loss: 219.5136 - val\_mae: 11.8986 - val\_mse: 219.2648


Epoch 2/30

**210/210**  3s 14ms/step - loss: 208.4257 - mae: 11.1222 - mse: 208.4255 - val\_loss: 202.5453 - val\_mae: 11.0873 - val\_mse: 202.2113


Epoch 3/30

**210/210**  3s 13ms/step - loss: 202.7528 - mae: 10.9684 - mse: 202.7528 - val\_loss: 202.1935 - val\_mae: 11.1603 - val\_mse: 201.9666


Epoch 4/30

**210/210**  3s 13ms/step - loss: 200.8208 - mae: 10.8961 - mse: 200.8214 - val\_loss: 199.8381 - val\_mae: 10.7330 - val\_mse: 199.5218


Epoch 5/30

**210/210**  3s 13ms/step - loss: 197.0804 - mae: 10.7782 - mse: 197.0802 - val\_loss: 201.0477 - val\_mae: 10.6818 - val\_mse: 200.7270


Epoch 6/30

**210/210**  3s 12ms/step - loss: 198.9585 - mae: 10.8274 - mse: 198.9591 - val\_loss: 197.6735 - val\_mae: 10.6539 - val\_mse: 197.4005


Epoch 7/30

**210/210**  3s 12ms/step - loss: 198.5806 - mae: 10.7971 - mse: 198.5815 - val\_loss: 196.6758 - val\_mae: 10.8178 - val\_mse: 196.4910


Epoch 8/30

**210/210**  3s 13ms/step - loss: 194.5566 - mae: 10.7437 - mse: 194.5569 - val\_loss: 197.7736 - val\_mae: 10.7849 - val\_mse: 197.4755


Epoch 9/30

**210/210**  3s 13ms/step - loss: 196.7931 - mae: 10.7770 - mse: 196.7944 - val\_loss: 204.1001 - val\_mae: 10.6567 - val\_mse: 203.7290


Epoch 10/30

**210/210**  3s 13ms/step - loss: 198.3773 - mae: 10.8183 - mse: 198.3757 - val\_loss: 201.8943 - val\_mae: 11.2187 - val\_mse: 201.7478


Epoch 11/30

**210/210**  3s 13ms/step - loss: 196.2389 - mae: 10.7865 - mse: 196.2393 - val\_loss: 195.8507 - val\_mae: 10.8117 - val\_mse: 195.6897


Epoch 12/30

**210/210**  3s 14ms/step - loss: 196.2606 - mae: 10.7662 - mse: 196.2606 - val\_loss: 197.2385 - val\_mae: 10.8333 - val\_mse: 197.0347


Epoch 13/30

**210/210**  3s 14ms/step - loss: 194.5317 - mae: 10.7362 - mse: 194.5317 - val\_loss: 196.3446 - val\_mae: 10.9340 - val\_mse: 196.2122


Epoch 14/30

**210/210**  3s 14ms/step - loss: 191.7983 - mae: 10.6656 - mse: 191.7988 - val\_loss: 196.6102 - val\_mae: 10.6663 - val\_mse: 196.4121


Epoch 15/30

**210/210**  3s 12ms/step - loss: 196.8160 - mae: 10.7610 - mse: 196.8164 - val\_loss: 196.0317 - val\_mae: 10.8560 - val\_mse: 195.7986


Epoch 16/30

**210/210**  3s 12ms/step - loss: 197.4550 - mae: 10.7620 - mse: 197.4549 - val\_loss: 194.4084 - val\_mae: 10.6706 - val\_mse: 194.1733


Epoch 17/30

**210/210**  3s 14ms/step - loss: 192.2394 - mae: 10.6451 - mse: 192.2399 - val\_loss: 193.0311 - val\_mae: 10.8362 - val\_mse: 192.9578


Epoch 18/30

**210/210**  3s 13ms/step - loss: 190.3562 - mae: 10.5848 - mse: 190.3559 - val\_loss: 192.2704 - val\_mae: 10.5375 - val\_mse: 192.2027


Epoch 19/30

**210/210**  3s 13ms/step - loss: 191.5047 - mae: 10.6351 - mse: 191.5045 - val\_loss: 192.2225 - val\_mae: 10.5080 - val\_mse: 192.1069

Epoch 20/30

**210/210**  3s 13ms/step - loss: 192.4837 - mae: 10.6682 - mse: 192.4836 - val\_loss: 198.8735 - val\_mae: 10.5644 - val\_mse: 198.6196

Epoch 21/30

**210/210**  3s 14ms/step - loss: 190.0445 - mae: 10.5599 - mse: 190.0448 - val\_loss: 194.7332 - val\_mae: 10.4838 - val\_mse: 194.6393

Epoch 22/30

210/210 ————— 3s 14ms/step - loss: 191.4072 - mae: 10.6043 - mse: 191.4075 - val\_loss: 207.6593 - val\_mae: 10.7024 - val\_mse: 207.4223  
Epoch 23/30

210/210 ————— 3s 14ms/step - loss: 195.8305 - mae: 10.6910 - mse: 195.8312 - val\_loss: 192.6539 - val\_mae: 10.8769 - val\_mse: 192.7865  
Epoch 24/30

210/210 ————— 3s 15ms/step - loss: 187.9753 - mae: 10.5507 - mse: 187.9752 - val\_loss: 190.0480 - val\_mae: 10.6299 - val\_mse: 190.0943  
Epoch 25/30

210/210 ————— 3s 13ms/step - loss: 187.6799 - mae: 10.5330 - mse: 187.6788 - val\_loss: 191.7179 - val\_mae: 10.4794 - val\_mse: 191.6578  
Epoch 26/30

210/210 ————— 3s 14ms/step - loss: 189.6170 - mae: 10.5656 - mse: 189.6173 - val\_loss: 190.8470 - val\_mae: 10.6809 - val\_mse: 190.8350  
Epoch 27/30

210/210 ————— 3s 15ms/step - loss: 189.3345 - mae: 10.5635 - mse: 189.3333 - val\_loss: 196.1583 - val\_mae: 10.9565 - val\_mse: 196.2539  
Epoch 28/30

210/210 ————— 3s 15ms/step - loss: 190.7337 - mae: 10.6078 - mse: 190.7346 - val\_loss: 193.9751 - val\_mae: 10.4601 - val\_mse: 193.8635  
Epoch 29/30


210/210 ————— 4s 18ms/step - loss: 186.8568 - mae: 10.4933 - mse: 186.8566 - val\_loss: 189.0742 - val\_mae: 10.5091 - val\_mse: 189.0633  
Epoch 30/30

210/210 ————— 3s 14ms/step - loss: 187.7057 - mae: 10.5109 - mse: 187.7060 - val\_loss: 190.4315 - val\_mae: 10.7709 - val\_mse: 190.5894

## Comparing losses with epochs

In [271...

```
def plot_history(history, key):
    plt.plot(history.history[key])
    plt.plot(history.history['val_'+key])
    plt.xlabel("Epochs")
    plt.ylabel(key)
    plt.legend([key, 'val_'+key])
    plt.show()
# Plot the history
plot_history(history, 'loss')
```



```
In [272...] z= model.predict(X_test)
```

1047/1047 ————— 1s 1ms/step

```
In [273...] r2_score(y_test, z)
```

```
Out[273]: 0.25638306639426456
```

## MAE RMSE MSE values for Neural Networks

```
In [275...] mse = mean_squared_error(y_test, z)
rmse = mse**.5
print("mse : ",mse)
print("rmse : ",rmse)
mae = mean_absolute_error(y_test, z)
print("mae : ",mae)
```

```
mse : 192.6221972548777
rmse : 13.878839910269075
mae : 10.842252765538582
```

```
In [ ]:
```